

Fluent, Composable Error Handling

Brian Beckman

August 12, 2014

Contents

1	Introduction	1
2	Motivating Example Problem	2
2.1	Fluent Solution in Java, No Error Handling	2
2.1.1	A Complete Sample Program	4
2.2	Fluent Solution in Java, with Exceptions	5
2.3	Fluency Lost Without Exceptions	7
3	Let's Do Better	12
3.1	Fluent Functional Solution With Exceptions	14
3.2	Fluent Error Handling Without Exceptions	16
3.2.1	Monads Are Values In Boxes	18
3.2.2	Monad Particulars and Generalities	20
3.2.3	Varying from the Standard Maybe Monad	20
4	Conclusion	20

1 Introduction

Consider a program composed of *serially dependent computations*, any of which produces either a value to send down the line, or an error. If any computation produces an error, no downstream computations should be attempted; the error should be the result of the entire sequence.

We show a sequence of programs in Java and Clojure for this scenario. We strive for fluent style, which minimizes the number of temporary variables to name and manage. This style directly mimics the abstract data flow of the solution.

We find that fluent style is only available in Java if errors are propagated by exception. Because the code for handling errors is in an independent *catch* block, the non-error code path can be written without explicitly handling errors.

With functional programming in general and Clojure in particular, we can have fluent code without exceptions because we can delegate error handling to a monad.

2 Motivating Example Problem

As a concrete example, suppose we get an authorization token, do a database lookup, do a web-service call, filter the results of that call, do another web-service call, and then combine the results. The data flow of our program resembles that in figure 1.

2.1 Fluent Solution in Java, No Error Handling

In the C++-like languages, including JavaScript and Java, we might keep intermediate results in instance variables and model the flow in object-to-object methods, that is, as **transforms**. This style is called **fluent style** because the text of the program resembles the flow in the diagram.

Imagine a *main* program like the following, in which each transform is on its own line, indented from its *source* by one 4-space tab stop. The **source** of a transform is an expression that produces an object to feed downstream.

```
public static void main(String[] args) {  
    Computation databaseResults = new Computation()  
        .authorize()  
        .readDatabase();  
    String result = databaseResults  
        .callWebService()  
        .filterResults()  
        .combineResults(databaseResults  
            .callOtherWebService());  
    System.out.println(result); }
```

The correspondence between the diagram of the program and the code is obvious. The code *looks like* the diagram. Given a change to the diagrammatic specification, a programmer may propagate the change straightforwardly to the code due to fluent style.

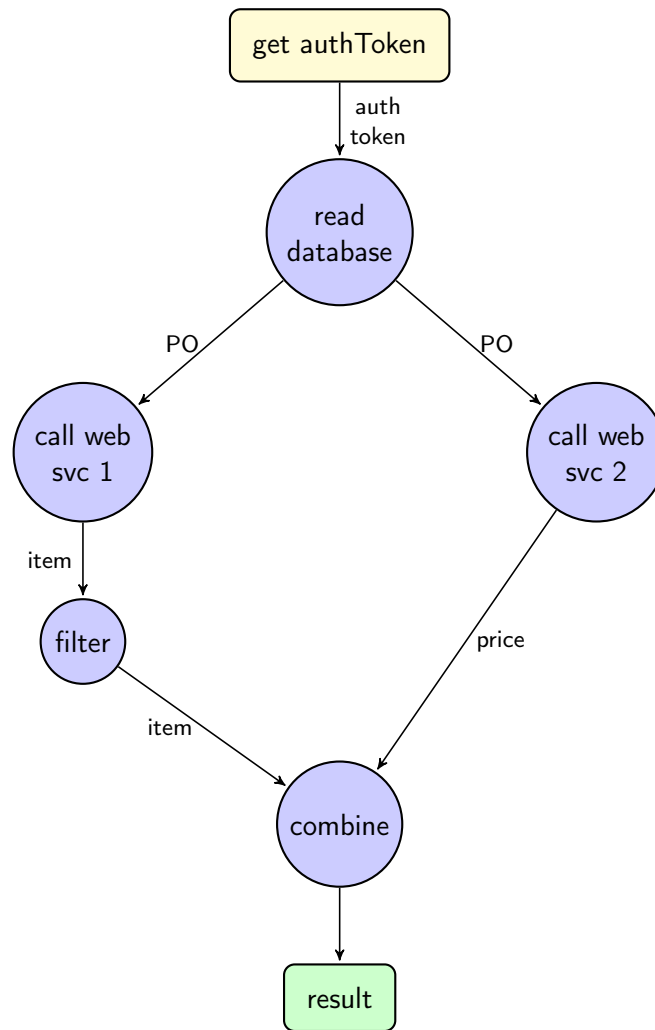


Figure 1: Serially dependent computations

We save the *Computation* produced by reading the database in its own local variable, namely *databaseResults*. The dataflow branches from that result; we need it twice, once for calling the first web service and once for calling the second web service. If not for this branching and recombining of the dataflow, we might have written the entire program as one, fluent expression with no intermediate variables.

Also note the non-idiomatic compressed style, minimizing blank lines and lines with just one closing brace. We adopt this style to save space in this paper; production versions of such code would have more white space.

2.1.1 A Complete Sample Program

The following is a complete program that mocks out the database and web-service results as static JSON objects encoded in strings. This can be compiled and executed, even online in a sandbox like http://www.compileonline.com/compile_java_online.php.

The most important thing to note about this code is that each method, *e.g.*, *authorize*, *readDatabase*, etc., takes an object – implicitly as *this* – and returns an object. This convention enables the fluent style by chaining transforms with dot. It so happens that we return *this* from each method. We could create a new object in each transform, but it would not simplify the code and would not give any performance advantage (quite the opposite). Instead, we propagate values through instance variables, namely *authToken*, *databaseResults*, etc.

```
public class Computation {
    private String authToken;
    private String databaseResults;
    private String webServiceCallResults;
    private String filteredWebServiceCallResults;
    private String otherWebServiceCallResults;

    public Computation () {}
    public Computation authorize() {
        authToken = "John's credentials";
        return this; }
    public Computation readDatabase() {
        databaseResults = "{\"name\":\"John\", \"PO\":\"421357\"}";
        return this; }
    public Computation callWebService() {
```

```

        webServiceCallResults =
            "[{\"item\":\"camera\"}, {\"item\":\"shoes\"}]";
        return this; }
    public Computation filterResults() {
        filteredWebServiceCallResults =
            "[{\"item\":\"camera\"}]";
        return this; }
    public Computation callOtherWebService() {
        otherWebServiceCallResults = "{\"price\":\"420.00\"}";
        return this; }
    public String combineResults(Computation other) {
        return "[" + filteredWebServiceCallResults +
            "," + otherWebServiceCallResults + "]]"; }

    public static void main(String[] args) {
        Computation databaseResults = new Computation()
            .authorize()
            .readDatabase();
        String result = databaseResults
            .callWebService()
            .filterResults()
            .combineResults(databaseResults
                .callOtherWebService());
        System.out.println(result);
    } }

```

2.2 Fluent Solution in Java, with Exceptions

The program above has *no* error handling. We *must* have error handling in real-world programs.

One of the better techniques for error handling in fluent style is with exceptions. If each sub-computation is responsible for throwing its own exception, then a single try-catch suffices to get error details out of the overall sequence, leaving the essential dataflow unchanged. Our main routine has minimal changes, and becomes simply

```

public static void main(String[] args) {
    try {
        Computation databaseResults = new Computation()
            .authorize()

```

```

        .readDatabase();
String result = databaseResults
        .callWebService()
        .filterResults()
        .combineResults(databaseResults
        .callOtherWebService());
System.out.println(result); }
catch (Exception e) {
    System.out.println(e.getMessage());
} }

```

Note, in passing, that we ignore resource management (database connections, sockets, file handles, etc.) in this paper.¹

Let us give each mocked sub-computation a 10% chance of erroring, and our entire sample becomes the following:

```

import java.util.Random;
public class Computation {
    private String authToken;
    private String databaseResults;
    private String webServiceCallResults;
    private String filteredWebServiceCallResults;
    private String otherWebServiceCallResults;
    private static Random random = new java.util.Random();
    private static Boolean randomlyError() {
        return random.nextDouble() < 0.10; }

    public Computation () {}
    public Computation authorize() throws Exception {
        if (randomlyError()) { throw new Exception("auth errored"); }
        authToken = "John's credentials";
        return this; }
    public Computation readDatabase() throws Exception {
        if (randomlyError()) { throw new Exception("database errored"); }
        databaseResults = "{\"name\":\"John\", \"PO\":\"421357\"}";
        return this; }
    public Computation callWebService() throws Exception {
        if (randomlyError()) { throw new Exception("ws1 errored"); }

```

¹Idiomatically, resources can be handled in a *finally* clause or with Java 7's Automatic Resource Management (ARM). See <http://bit.ly/15GYkMh>

```

        webServiceCallResults =
            "[{\"item\":\"camera\"}, {\"item\":\"shoes\"}]";
        return this; }
    public Computation filterResults() throws Exception {
        if (randomlyError()) { throw new Exception("filter errored"); }
        filteredWebServiceCallResults =
            "[{\"item\":\"camera\"}]";
        return this; }
    public Computation callOtherWebService() throws Exception {
        if (randomlyError()) { throw new Exception("ws2 errored"); }
        otherWebServiceCallResults = "{\"price\":\"420.00\"}";
        return this; }
    public String combineResults(Computation other) throws Exception {
        if (randomlyError()) { throw new Exception("combine errored"); }
        return "[" + filteredWebServiceCallResults +
            "," + otherWebServiceCallResults + "]]"; }

    public static void main(String[] args) {
        try {
            Computation databaseResults = new Computation()
                .authorize()
                .readDatabase();
            String result = databaseResults
                .callWebService()
                .filterResults()
                .combineResults(databaseResults
                    .callOtherWebService());
            System.out.println(result); }
        catch (Exception e) {
            System.out.println(e.getMessage());
        }
    }
}

```

2.3 Fluency Lost Without Exceptions

Error handling with exceptions is debatable,² especially in Java, where runtime exceptions need not be declared,³ but the alternative of checked exceptions can be considered harmful.⁴

²<http://www.joelonsoftware.com/items/2003/10/13.html>

³<http://bit.ly/1e5P6Cg>

⁴<http://bit.ly/9NyrD>

Worse yet, the semantics of composed locks and exceptions are black magic. The fundamental reason is that an exception thrown from inside a lock leaves the program in an indeterminate state for other threads, with the lock summarily abandoned. There are expert techniques for mitigating this,⁵ but a defensible way out is just to eschew exceptions.

But, rather than join the debate, just imagine that we have decided against exceptions for whatever reason and try to write reasonable code.

Add a private *String* field *errorResult*, and let every method set the error result if and only if it errors. We must change *combineResults*; it can no longer return just a *String*, but rather a *Computation*, because it may, itself, produce an error. Furthermore, we lose the fluent style because every call must be individually checked.

A particularly nasty way to check every call is as follows:

```
public static String computation () {
    Computation c1 = new Computation();
    Computation c2 = c1.authorize();
    if (c2.errorResult.isEmpty()) {
        Computation c3 = c2.readDatabase();
        if (c3.errorResult.isEmpty()) {
            Computation c4 = c3.callWebService();
            if (c4.errorResult.isEmpty()) {
                Computation c5 = c4.filterResults();
                if (c5.errorResult.isEmpty()) {
                    Computation c6 = c3.callOtherWebService();
                    if (c6.errorResult.isEmpty()) {
                        Computation c7 = c5.combineResults(c6);
                        if (c7.errorResult.isEmpty()) {
                            return c7.getResult(); }
                        else {return c7.errorResult;} }
                    else {return c6.errorResult;} }
                else {return c5.errorResult;} }
            else {return c4.errorResult;} }
        else {return c3.errorResult;} }
    else {return c2.errorResult;} }
public static void main(String[] args) {
    System.out.println(computation()); }
}
```

This is so intolerable as to barely deserve criticism, despite the fact that

⁵<http://bit.ly/q001r>

its working set is optimized for the positive path!⁶ We've lost any correspondence between the program text and the program specification, *i.e.*, the diagram in figure 1. All options for nesting and placement of curly braces are ludicrous. Changing the computation graph would entail a sickening amount of work. Code like this is best left to automatic code generators, if we tolerate it at all.

The prevailing style, nowadays, is to reverse error branches and to return as early as possible from the main routine. I have reviewed many instances of this style in shipped code from pre-eminent shops. Multiple returns were condemned in the dogma of structured programming. They are also lethal in code that manages resources,⁷. Despite these drawbacks, justification for multiple returns is three-fold:

- it results in linear code that can be read from top to bottom
- edits to the computation graph entail just adding or subtracting a localized block of a few lines of code and adjusting a few temporary variables
- modern compilers can reverse the branches *again* in the generated code automatically after profiling⁸

This alternative⁹ is the following:

```
public static String computation() {
    Computation c1 = new Computation();
    Computation c2 = c1.authorize();
    if (! c2.errorResult.isEmpty()) {return c2.errorResult;}
    Computation c3 = c2.readDatabase();
    if (! c3.errorResult.isEmpty()) {return c3.errorResult;}
    Computation c4 = c3.callWebService();
    if (! c4.errorResult.isEmpty()) {return c4.errorResult;}
    Computation c5 = c4.filterResults();
    if (! c5.errorResult.isEmpty()) {return c5.errorResult;}
    Computation c6 = c3.callOtherWebService();
    if (! c6.errorResult.isEmpty()) {return c6.errorResult;}
    Computation c7 = c5.combineResults(c6);
```

⁶The error branches are all at addresses far from the common-case, non-error branches, which are clustered together for maximum locality.

⁷<http://bit.ly/sAvDmY>

⁸<http://bit.ly/QkXSM>

⁹favored in the previously cited Joel-on-Software blog

```

        if (! c7.errorResult.isEmpty()) {return c7.errorResult;}
        return c7.getResult(); }
public static void main(String[] args) {
    System.out.println(computation()); }

```

This, at least, gets rid of the ludicrous nesting, but exposes another deep weakness: a proliferation of temporary variables to hold the intermediate *Computations*. Why bother with this when we have no hope of fluent style? Instead, consider

```

public static String computation() {
    Computation c1 = new Computation();
    c1.authorize();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.readDatabase();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.callWebService();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.filterResults();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.callOtherWebService();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.combineResults(c1);
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    return c1.getResult(); }
public static void main(String[] args) {
    System.out.println(computation()); }

```

Edits to the graph now entail even easier edits to the source.

The whole program at this point is the following:

```

import java.util.Random;
public class Computation {
    private String errorResult;
    private String result;
    private String authToken;
    private String databaseResults;
    private String webServiceCallResults;
    private String filteredWebServiceCallResults;
    private String otherWebServiceCallResults;
    private static Random random = new java.util.Random();

```

```

private static Boolean randomlyError() {
    return random.nextDouble() < 0.10; }

public Computation () {errorResult=""; result="no result";}
public Computation authorize() {
    if (randomlyError()) { errorResult = "auth errored"; }
    authToken = "John's credentials";
    return this; }
public Computation readDatabase() {
    if (randomlyError()) { errorResult = "database errored"; }
    databaseResults = "{\"name\":\"John\", \"PO\":\"421357\"}";
    return this; }
public Computation callWebService() {
    if (randomlyError()) { errorResult = "ws1 errored"; }
    webServiceCallResults =
        "[{\"item\":\"camera\"}, {\"item\":\"shoes\"}]";
    return this; }
public Computation filterResults() {
    if (randomlyError()) { errorResult = "filter errored"; }
    filteredWebServiceCallResults =
        "[{\"item\":\"camera\"}]";
    return this; }
public Computation callOtherWebService() {
    if (randomlyError()) { errorResult = "ws2 errored"; }
    otherWebServiceCallResults = "{\"price\":\"420.00\"}";
    return this; }
public Computation combineResults(Computation other) {
    if (randomlyError()) { errorResult = "combine errored"; }
    result = "[" + filteredWebServiceCallResults +
        "," + otherWebServiceCallResults + "]";
    return this;}
public String getResult() {return result;}
public static String computation() {
    Computation c1 = new Computation();
    c1.authorize();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.readDatabase();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
    c1.callWebService();
    if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
}

```

```

        c1.filterResults();
        if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
        c1.callOtherWebService();
        if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
        c1.combineResults(c1);
        if (! c1.errorResult.isEmpty()) {return c1.errorResult;}
        return c1.getResult(); }
    public static void main(String[] args) {
        System.out.println(computation());
    }
}

```

3 Let's Do Better

Looking back, the main benefits of fluent style are

- direct correspondence between the program specification and the program text – the text *looks like* the diagram
- edits to the specification and edits to the code are straightforward and parallel
- minimal number of temporary variables

But we only have fluent style if we use exceptions. Without exceptions, we're essentially writing assembly language: storing and combining intermediate results in temporary variables and checking for errors after every step.

It's possible to do much better, with and without exceptions, by going *functional*.

In Java, the fundamental modeling tool is the *mutable, stateful object*. Stateful-object programming has many disadvantages:

- dataflow is awkward to model
- concurrency requires locks, which are complex; and *global reasoning*, which is very complex, to avoid deadlock, livelock, cycles, starvation, priority inversion, etc.
- concurrency with exceptions is very difficult
- composing stateful objects, even without concurrency, is difficult: the operational semantics of even a sequential program requires *temporal*

reasoning, well outside the capabilities of compilers and programming tools

To begin a trek away from stateful-object programming, start with a question. Why did we use a new instance variable in the object, *e.g.*, *authToken*, *databaseResults*, *webServiceCallResults*, etc., for each intermediate state? Why not reuse a single variable for all non-error results? After all, we used a single variable for the error result?

The reason is that we wanted the individual methods that update non-error state to be as independent as possible. Though our mocks don't do so, in a real program, each intermediate computation would use the result of its predecessors: *readDatabase* would use the *authToken*, the web-service calls would use *databaseResults* and so on. With a separate, named variable for each intermediate result, the correctness of individual sub-computations would be easier to verify by inspection. With a single, re-used *String* variable, temporal flow would be even more obscured and the program would be even more difficult to understand and maintain. It is definitely worth a few more named variables to make the program easier for the next programmer tasked with reading the code. Because the only tool in Java is mutable state, it is hard to do better than a sequence of instance variables mirroring the sequence of sub-computations.

The essence of the problem is in modeling a *flow* of data through *transforms* as a flow of data through mutable variables. If, instead, we invert the paradigm to put the *transforms* under focus, we sidestep this problem. Doing so requires a language with first-class transforms, that is, *functions*. Mutable state variables become immutable function parameters. Thread-safety becomes automatic and locks do not arise.

C# has first-class functions, a.k.a. *lambda expressions*, as does C++ 11 and as will Java 8. In the mean time, we can use *Clojure*, a Java-compatible functional language.

As to error handling, fluency is immediate with exceptions, as before. We achieve fluency for errors-as-return-values with a *monad*. Despite the name, monads are not exotic or complicated. They furnish a straight generalization of function composition. Our monad flows values through the pipeline, applying user-supplied functions only when an error has not been produced. The user-supplied functions do not need to check. This monad abstracts the boilerplate code for checking errors.

Other monads abstract other boilerplate into the monad machinery, leaving user-supplied code free of clutter. Monads are inherently stateless, so are not a good fit in stateful-object programming languages. They are funda-

mental for concurrent and distributed programming, and their benefits are evident in functional programming languages.

3.1 Fluent Functional Solution With Exceptions

We may write the program with only one intermediate variable, as before. This variable holds results of the database read, which we must use for each of the two web-service calls. There are many techniques for eliminating this variable, such as *memoization*, *common sub-expression elimination*, *lambda lifting*, *the state monad*, or *parallel composition*. That is for another time and place. For now, write the flow directly as a sequential composition of function calls *via* Clojure's `->`, using its `let` syntax for the one intermediate variable, as follows:

```
(try
  (let [db-results ; be the result of the composition
        (-> (computation)
              authorize
              read-database
              )]
    (-> db-results
        call-web-service
        filter-ws
        (combine (-> db-results
                     call-other-web-service))))
  (catch Exception e (.getMessage e)))
```

This looks very much like the fluent Java solution. In Java, we have fluent streams of items connected by dots. In Clojure, we have the same fluent streams of items headed by arrows. The fundamental difference is that Clojure has no instance variables, therefore it is automatically re-entrant, unlike our Java solution. The semantics of the Clojure program is evident in the text – it's a flow of data through serially dependent computations. There is no temporal component to the semantics. Contrast with the Java solution, where we must understand the temporal sequence of operations to understand the program.

The rest of the Clojure program is as follows. First, declare a namespace for our symbols to inhabit:

```
(ns temp-1.core)
```

Define a private symbol, *computation*, to be a function of no arguments (denoted by the empty square braces) that produces an empty hash-map (denoted by the empty curly braces).

```
(defn- computation [] {})
```

Define a private symbol *randomly-error* to be a function of no arguments that produces *true* if a random double-precision number is less than 0.10.

```
(defn- randomly-error [] (< (rand) 0.10))
```

Define several private symbols to be mocks of functions that produce application-specific data, or, with 10% probability, throw an exception. These all use three-term *if*-branch forms, equivalent to Java's $\langle e_1 \rangle ? \langle e_2 \rangle : \langle e_3 \rangle$ operator, taking an expression of Boolean type and two arbitrary expressions, only one of which is evaluated.

```
(defn- authorize [computation]
  (if (randomly-error) (throw (Exception. "auth errored"))
      {:auth-token "John's credentials"}))
(defn- read-database [auth-token]
  (if (randomly-error) (throw (Exception. "database errored"))
      {:name "John", :PO 421357}))
(defn- call-web-service [database-results]
  (if (randomly-error) (throw (Exception. "ws1 errored"))
      [{:item "camera"}, {:item "shoes"}]))
(defn- filter-ws [web-service-call-results]
  (if (randomly-error) (throw (Exception. "filter errored"))
      [{:item "camera"}]))
(defn- call-other-web-service [database-results]
  (if (randomly-error) (throw (Exception. "ws2 errored"))
      [{:price 420.00M}]))
```

The last application-specific function, *combine* takes two arguments and combines them.

```
(defn- combine [filtered-web-service-results
                other-web-service-call-results]
  (if (randomly-error)
      (throw (Exception. "combine errored"))
      (concat filtered-web-service-results
                other-web-service-call-results)))
```

Several improvements are notable in this attempt:

- first, as stated, with only one exception, state variables have become immutable parameters to re-entrant functions, purely local to each transform
- the one remaining intermediate variable is itself immutable, removing any need for temporal reasoning – we only need to understand dependencies, and they are explicit in the code
- the code is shorter, less repetitive, less noisy
- the values of each mock can be modeled directly as hash-maps, arrays, integers, and decimal numbers like *420.00M*, as opposed to JSON objects encoded in strings
 - such direct modeling removes the implied need, unstated in our Java solution, for JSON-processing code
 - such direct modeling also means that we do not need direct Java interop; our *computation* “constructor” just returns an empty hash-map
 - if we did need need an existing Java class, we would only need to *import* the class and change our constructor call from `(computation)` to `(Computation.)`, shorthand for `(new Computation)`

We continue to use Java’s native *Exception* class.

3.2 Fluent Error Handling Without Exceptions

The improvements above alone justify the Clojure solution over the Java solution. But the case is really obvious when we get to error handling without exceptions. In Clojure, we use a variation of **the Maybe monad**.¹⁰

¹⁰<http://bit.ly/WV02FF>

Here is the main code:

```
(let [db-results
      (=>> (computation)
            authorize
            read-database)]
  (=>> db-results
    call-web-service
    filter-ws
    (combine (=>> db-results
                  call-other-web-service))))
```

It looks *just like* the non-monadic code, only with a different arrow. The prerequisite helpers are as follows:

```
(defn- computation [] (with-em-result {}))
(defn- authorize [computation]
  (with-em-result
    (if (randomly-error) {:error "auth errored"}
        {:auth-token "John's credentials"})))
(defn- read-database [auth-token]
  (with-em-result
    (if (randomly-error) {:error "database errored"}
        {:name "John", :PO 421357})))
(defn- call-web-service [database-results]
  (with-em-result
    (if (randomly-error) {:error "ws1 errored"}
        [{:item "camera"}, {:item "shoes"}])))
(defn- filter-ws [web-service-call-results]
  (with-em-result
    (if (randomly-error) {:error "filter errored"}
        [{:item "camera"}])))
(defn- call-other-web-service [database-results]
  (with-em-result
    (if (randomly-error) {:error "ws2 errored"}
        [{:price 420.00M}])))
```

All but *combine* are straightforward modifications of the non-monadic code, simply wrapping their results in a *with-em-result*. In fact, this wrapping

is one of only two things we must learn about monads: **it puts values in boxes.**

3.2.1 Monads Are Values In Boxes

An instance of a monad is just a value in a box. Every monad has an operator, *m-result*, which takes a value and puts it in a box. All monads work this way: take a value and put it in a box. The box can be arbitrarily complicated inside, and each type of monad has its own type of box. But all monads have the *m-result* operator in common.

The *with-em-result* macro used above is just shorthand *m-result* for our error-propagating monad. Here is its definition:

```
(defmacro with-em-result [expr]
  `(with-monad if-not-error-m (m-result ~expr)))
```

Monads have one more essential operator, *m-bind*. This takes two arguments: a value-in-a-box and a function-that-puts-a-value-in-a-box. This function's signature is just like the signature of *m-result*: it takes a value and puts it in a box. Here we see again the fundamental simplifying idea: put values in boxes.

Our monad is called *if-not-error-m*. Its *m-result* operator simply returns its input. This monad's box is an invisible box. That's ok; monadic boxes can be arbitrarily complicated, including not complicated at all. A real-world implementation might do logging as a side-effect. Our monad's *m-bind* takes a value-in-a-box and a user-supplied function, like the *m-binds* of all monads. With this monad, *m-bind* applies the user-supplied function only if the value doesn't have an error in it. Otherwise, it puts the value, including its error, in a box and passes it down the line. Since every subsequent element of the pipeline must go through the very same *m-bind*, the very first error produced in the pipeline is propagated all the way through. All subsequent user-supplied functions are sidestepped. If there is no error, it applies the user-supplied function to the value. The user-supplied function has the responsibility of putting the value in a box, automated by our *with-em-result* macro.

```
(defmonad if-not-error-m
  [m-result (fn [value] value)
  m-bind    (fn [value f]
               (if-not (:error value)
```

```

        (f value)
        (m-result value)))
  ])

```

We can see how our new arrow operator, `=>>`, composes values through *m-bind*. For instance, the fragment

```

(=>> (computation)
      authorize
      read-database)

```

expands approximately to

```

(with-monad if-not-error-m
  (m-bind
    (m-bind (computation)
              (fn [temp] (authorize temp))))
    (fn [temp] (read-database temp)))))

```

Longer chains expand to more deeply nested compositions of *m-bind* behind the scenes, precisely where it should be. Contrast this to the Java situation, where we had no tools in the language for mitigating intolerable nesting and proliferation of temporaries. Here, we define our arrow directly in our application code. Clojure, in particular, and homoiconic languages,¹¹ in general, often have embedded rewrite systems¹² or macro languages.¹³ Without going more deeply into Clojure’s macro syntax or the monadic *m-chain* helper it employs, the arrow macro is at least small:

```

(defmacro =>> [in-monad & transforms]
  `(with-monad if-not-error-m
     ((m-chain [~@transforms]) ~in-monad)))

```

The final monadic helper, *combine*, illustrates *m-bind* directly. Given a value from the second of the two web-service calls, it produces a *closure*¹⁴ a first-class value of type *function-that-puts-a-value-in-a-box*. That function will be composed, *via m-bind*, with the filtered value of the first of the two web-service calls. We use the suffix *-val* to distinguish values-not-in-boxes from value-in-boxes, which are produce by *with-em-result*.

¹¹<http://en.wikipedia.org/wiki/Homoiconicity>

¹²<http://en.wikipedia.org/wiki/Rewriting>

¹³[http://en.wikipedia.org/wiki/Macro_\(computer_science\)](http://en.wikipedia.org/wiki/Macro_(computer_science))

¹⁴[http://en.wikipedia.org/wiki/Closure_\(computer_science\)](http://en.wikipedia.org/wiki/Closure_(computer_science))

```
(defn- combine [other-ws-results-val]
  (fn [filtered-ws-results-val]
    (with-em-result ; produces a value-in-a-box
      (if (randomly-error)
        {:error "combine errored"}
        (concat filtered-ws-results-val
                  other-ws-results-val))))))
```

3.2.2 Monad Particulars and Generalities

The implementations of *m-result* and *m-bind* are particular to each monad. Package your application-dependent logic in user-supplied functions-that-put-values-in-a-box. Compose them *via* *m-bind*, *m-chain*, or the `=>>` operator.

3.2.3 Varying from the Standard Maybe Monad

Why do we have a variation of the standard *Maybe* monad? Why not just use *Maybe*? First, *Maybe* produces an unadorned *Nothing* if anything goes wrong. The consumer of the computation doesn't know what stage of the pipeline failed nor any details at all about the error. Such is not tolerable in the real world. In our example of a database read followed by web-service calls followed by filtering and combining, if something goes wrong in this sequence of computations, we need to know exactly where and to log as much detail as we can get about the failure. But we certainly don't want any computations downstream of the failure to be attempted. We want the pass-through semantics of *Maybe* but without the ignorance.

4 Conclusion

Controlling complexity is the central problem of software engineering. Making code closer to specifications is essential, and that is the central theme of progress in programming languages.

Many business processes are fundamentally data-flows. Modeling data-flows directly in our programming languages and applications is a clear advantage.

Fluent style is a direct representation of dataflow dependencies in the text of a program. Stateful-object programming is an awkward fit to fluent

style. The semantics of stateful-object programs are temporal and implicit, while the semantics of fluent style are dependency-based and explicit.

Fluent style can be approached in stateful-object programming by simulating flow-through-transforms as flow-through-instance-variables, but the style is brittle. For instance, changing from errors-as-exceptions to errors-as-return-codes makes fluent style intractable.

Furthermore, stateful-object programming makes concurrency more hazardous. Composability of objects and exceptions is difficult to test, and manual, pattern-based concurrent programming with locks is a dark art for the elite few. Yet, concurrency is necessary in everyday web programming.

In addition to enabling fluent style everywhere, functional programming alleviates the other hazards of stateful-object programming. It is composable both with exceptions and with errors-as-return-codes. Because function-parameters are immutable, functions are naturally re-entrant. Functional programming also sidesteps the temporal reasoning required for stateful-object programming, making code easier to test and debug.